
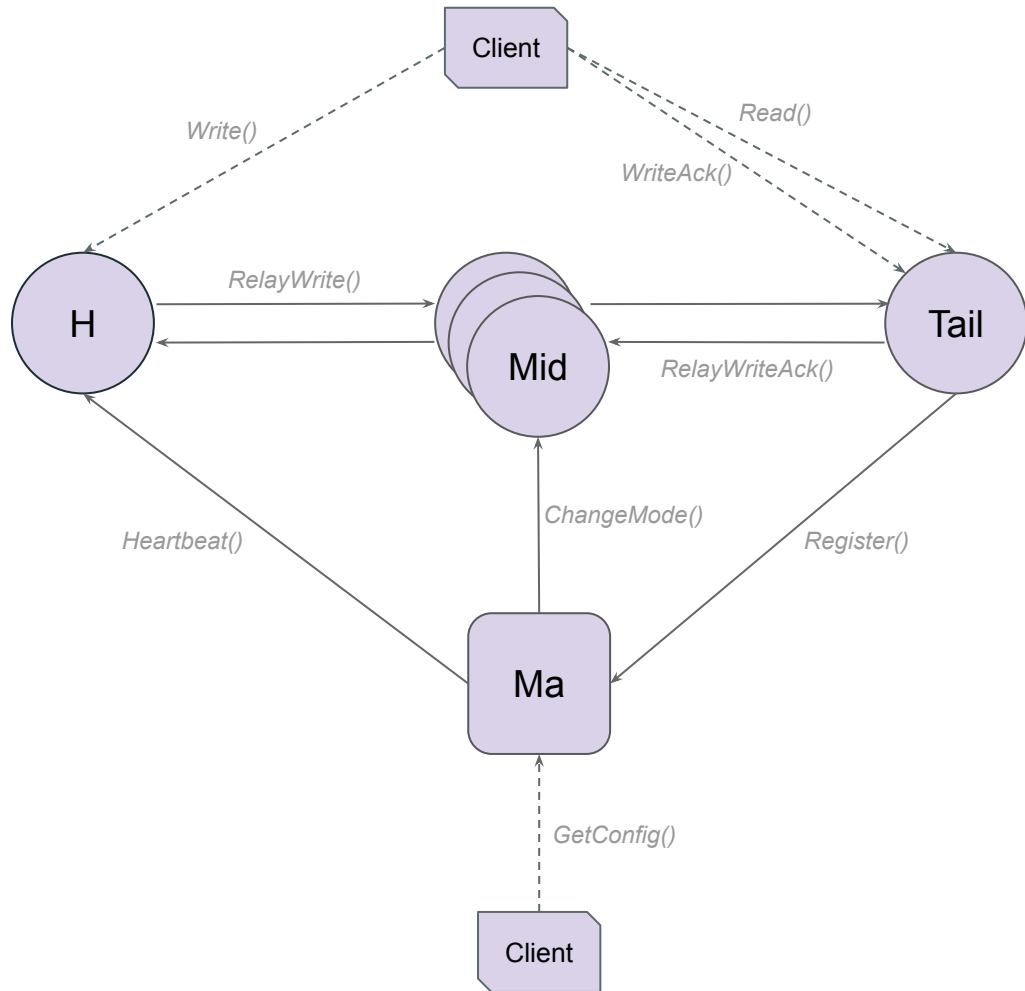


# *Chain Replicated Block Store*



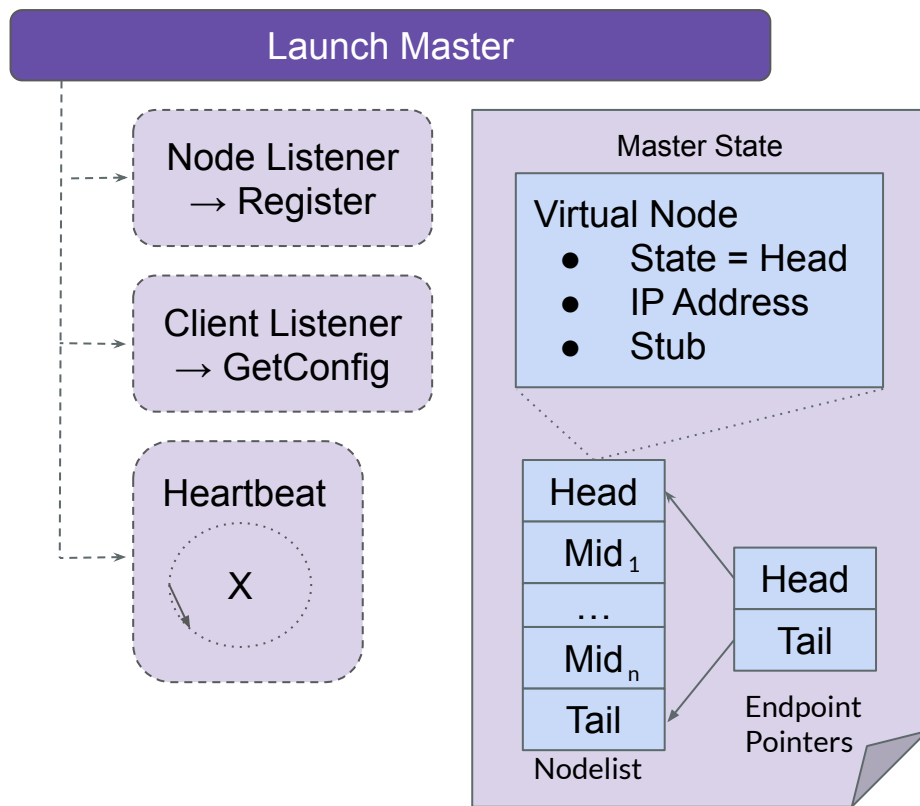
# Architecture

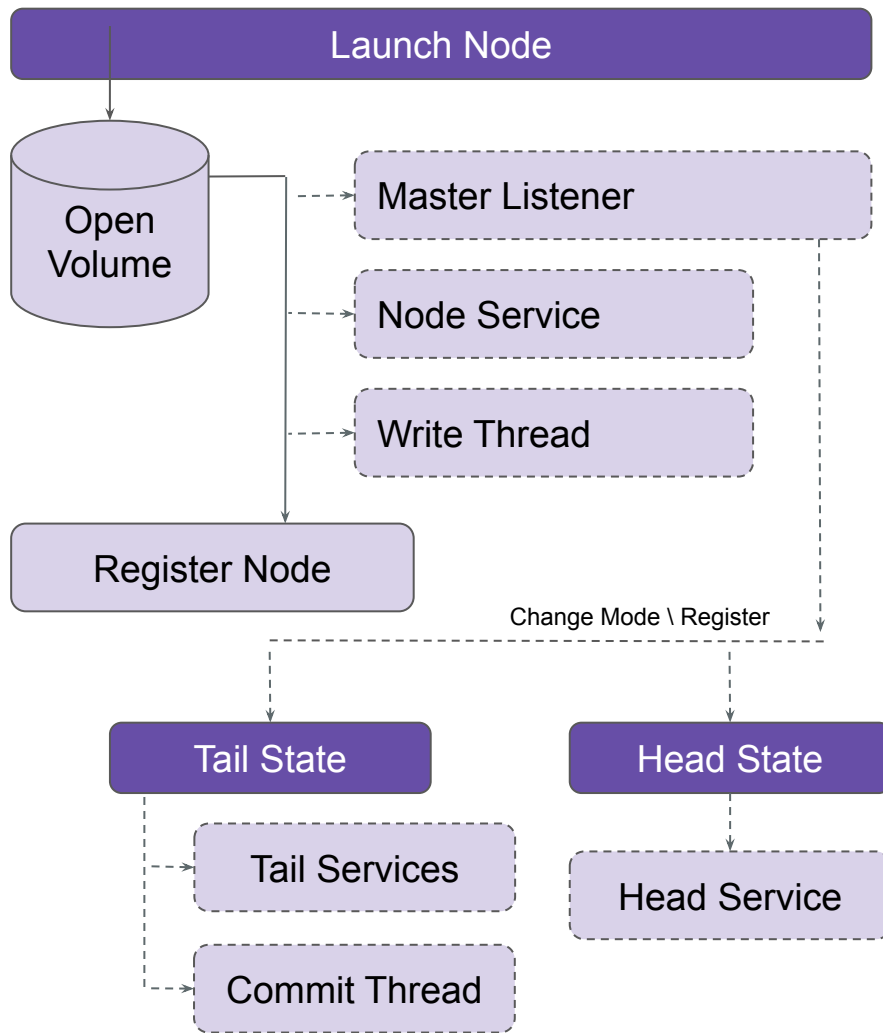
- Multi-Node Chain Replication
  - 1 to N nodes in chain
  - Chain managed by Master
  - Client gets chain information from master, speaks with head and tail
- Consistency
  - Linearizable up to N-1 Failures
  - Durable and Consistent up to N Failures
- C\C++ with gRPC
- All benchmarks run on cloud nodes
  - Ubuntu 20.04
  - 4Gb Memory
  - 10 core 2.6GHz cpu
  - HDD storage



# Architecture - Master

- Coordinates chain
  - Tracks nodes in chain
  - Handles membership
  - Informs nodes of state changes
- Communicates Endpoints to client





## Architecture - Node

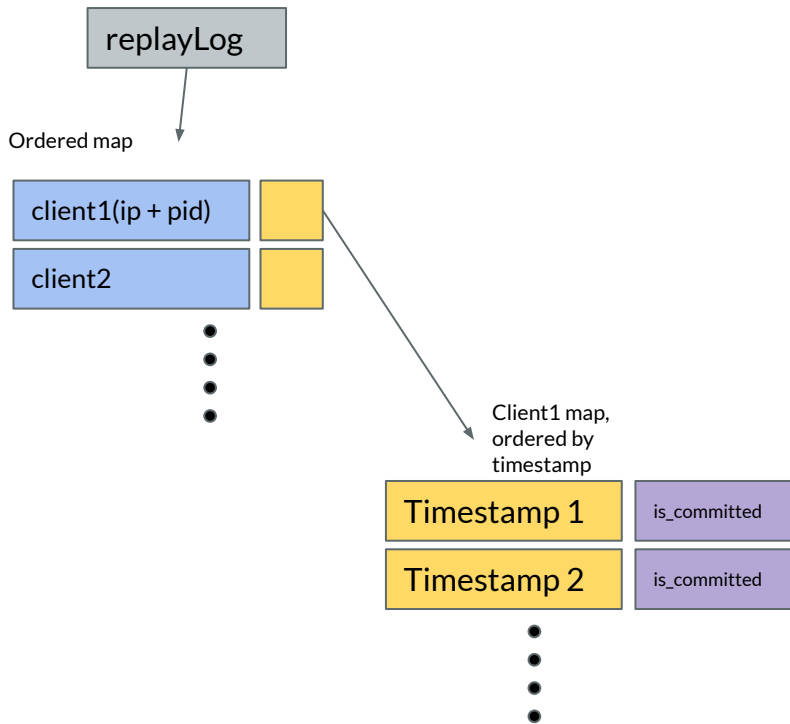
- Nodes run a variety of services that facilitate communication
- Master can alter state, which changes services exposed
- Blocks stored in a volume

# Client Library

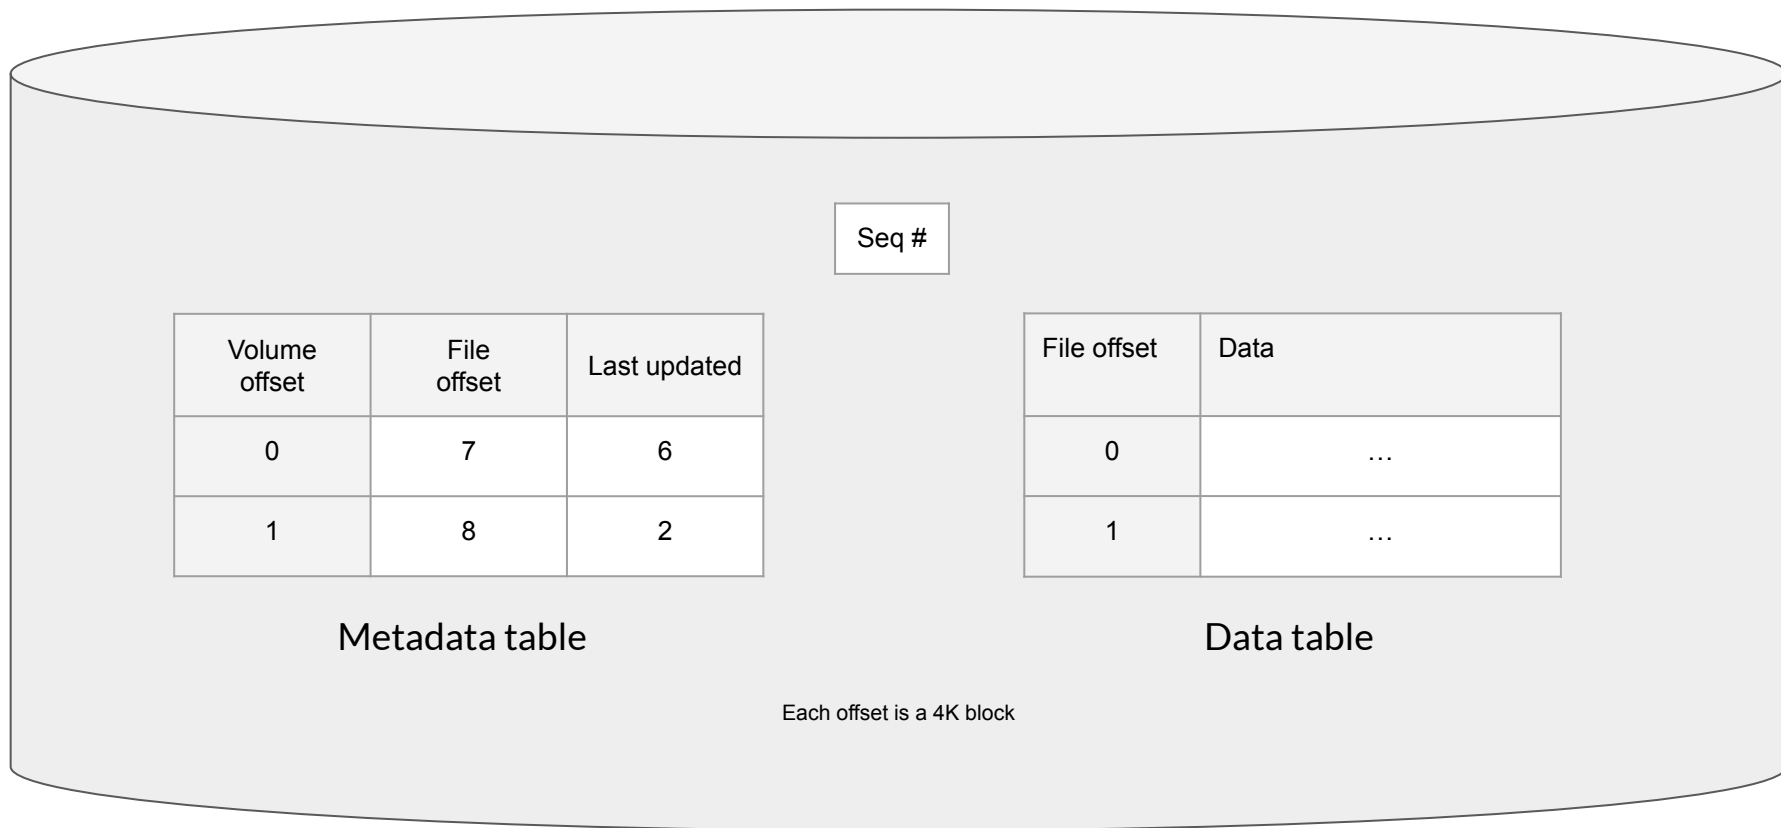
- Clients generate request IDs(clientIp: clientPid:timestamp) for each write
  - A pair of Client IP : Client PID is generated upon library class initialization
  - Timestamp is generated upon each write
- **Read (&data, file\_offset)**
- **Write (data, file\_offset, &timestamp)**
  - Generate the timestamp upon call
  - Return right away. Add the write to a list of pending writes
- **AckWrite(timestamp)**
  - Call tail to see if the write for a give timestamp is committed
  - This function also return as soon as tail gets the result. If the client wants to do things like retry after x seconds of write not acked, they will need to use their own timer.
- **RetryTopPendingWrite(&timestamp)**
  - Resend the request with a certain. This retry will only be sent if the timestamp is present in the list of pending writes
- **PopPendingWrite(&timestamp, &pending\_write\_entry) and PeekPendingWrite(&timestamp, &pending\_write\_entry)**
  - Both check top of the pending write list. Can be used to remove acked writes or skip writes that you don't want to acknowledge at all
- **RefreshConfig():**
  - Private method. Contacts the master and requests the address of the head and tail nodes. This is called whenever a request to a head or tail fails with connection error

# Replay Log

- **AddToLog(client\_request\_id)**
  - Called by every node: Adds incoming client request IDs to replay log
  - Used to detect retries. Return appropriate value based on whether the request id is already present in the log
  - Log added as non-committed
- **CommitLogEntry(client\_request\_id)**
  - Called by any node when the request has been committed(indicate by RelayWriteAck grpc call passed upstream)
- **AckLogEntry(client\_request\_id)**
  - Tail node calls this when client calls in ack. Removes a committed entry, and any entry belong to the same client that is older(assume client will try to ack request in order)
  - Any change to replay log is passed upstream
- **CleanOldLogEntry(age);**
  - Replay log garbage collects stale Client request ids that is too old



# Volume



# Volume

## Atomic commit

- Copy-on-write for data
- Copy-on-write for metadata
- Multi-level table for metadata
- Scales well with larger volumes

```
Write(data, vol_offset, seq_num):  
    write(data, new_file_offsets);  
    write(metadata, new_md_offsets);  
    record_uncommitted(vol_offset, new_file_offsets);  
    record_pending(seq_num, new_md_offsets);
```

```
Commit(seq_num, vol_offset):  
    get_pending(seq_num);  
    fsync();  
    write(seq_num, new_md_tab, 0);  
    fsync();  
    add_old_blocks_to_free_list();
```

## Non-atomic commit

- Copy-on-write for data
- Single level metadata table
- Faster metadata accesses
- May have partial writes if all nodes crash

```
Write(data, vol_offset, seq_num):  
    write(data, new_file_offset);  
    record_uncommitted(vol_offset, new_file_offset);
```

```
Commit(seq_num, vol_offset):  
    get_uncommitted(seq_num);  
    write_metadata(vol_offset, new_file_offset);  
    fsync();  
    write(seq_num, 0);  
    fsync();  
    add_old_blocks_to_free_list();
```






# Volume

	Atomic	Non-atomic
Aligned write	19 us	4.6 us
Aligned read	3 us	2.3 us
Unaligned write	25.6 us	10.5 us
Unaligned read	5.9 us	4.1 us
Aligned commit	1.7 ms	2.5 ms
Unaligned commit	1.7 ms	5.8 ms
Write-commit	12 ms	12 ms

# Message Propagation


- Head Service
  - Write() adds to pending list and replay log
- Node Service
  - Background 'write thread' pulls from pending list, writes, forwards, and adds to sent list
  - RelayWrite() adds to pending list and replay log
  - RelayWriteAck() removes from sent list, commits, and forwards
  - AckReplayLog() removes from replay log and forwards
- Tail Service
  - Background 'commit thread' pulls from sent list, commits, and forwards
  - Read() returns data
  - WriteAck() checks commit, removes from replay log, and forwards

# Message Propagation

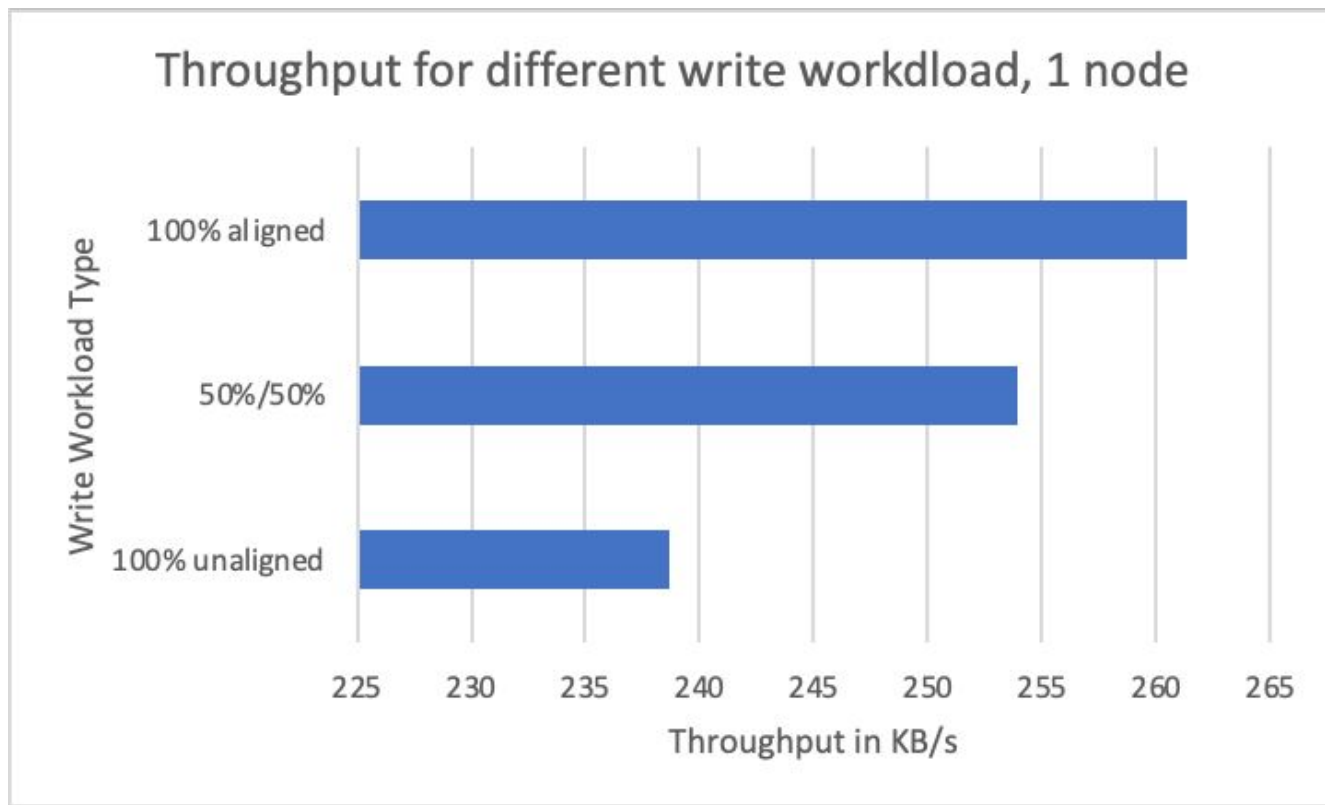


# Latency benchmark - write


Latency with different number of nodes



# Throughput benchmark - write mode



# Throughput benchmark - effect of node number



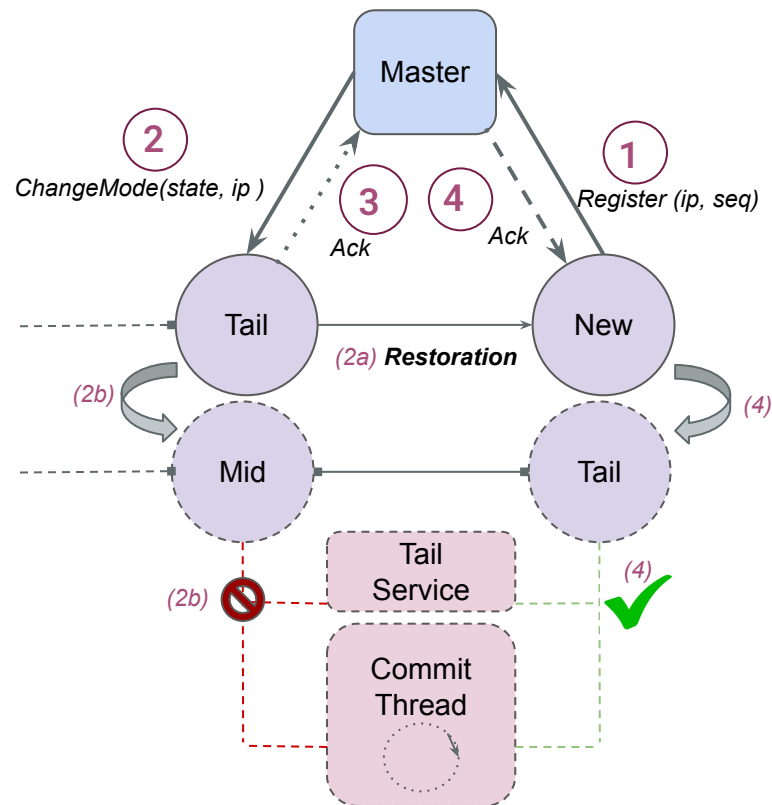

# Registration & Restoration

## Existing Chain

- 1) New nodes register with master, master generates virtual node
- 2) Master notifies tail of new state, tail generates downstream channel
  - a) Tail restores new node
  - b) updates own state, kills tail service
- 3) Master adds new node to node list and updates tail pointer
- 4) New node updates internal state, launches tail service and commit thread

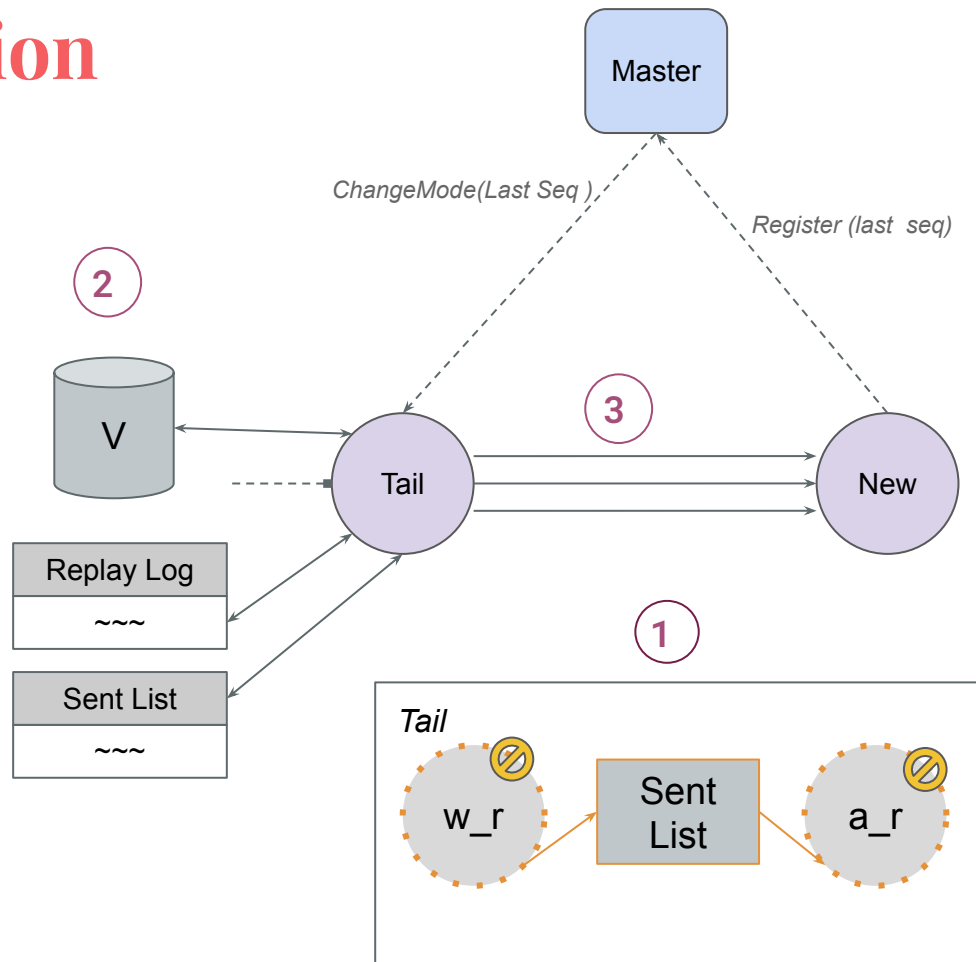
## New Chain

- Skip step 2
- New node becomes Single Server
- Launches Head & Tail services \ commit thread




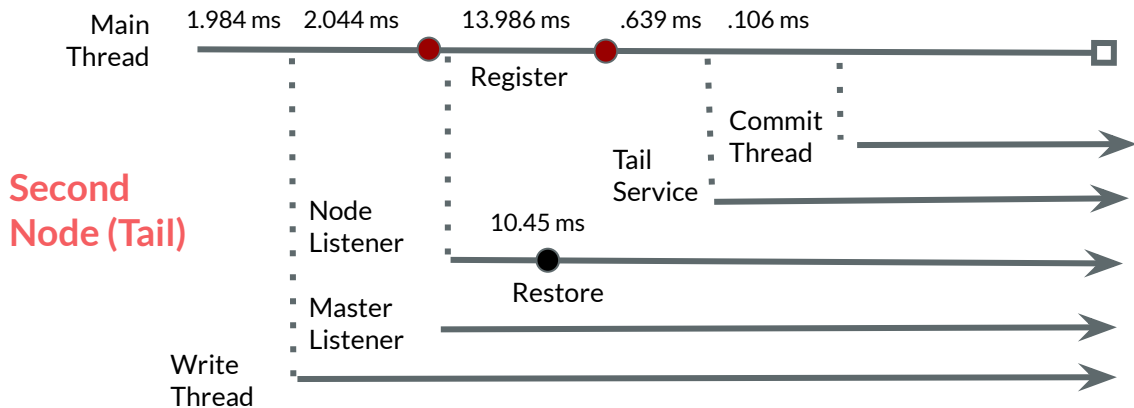
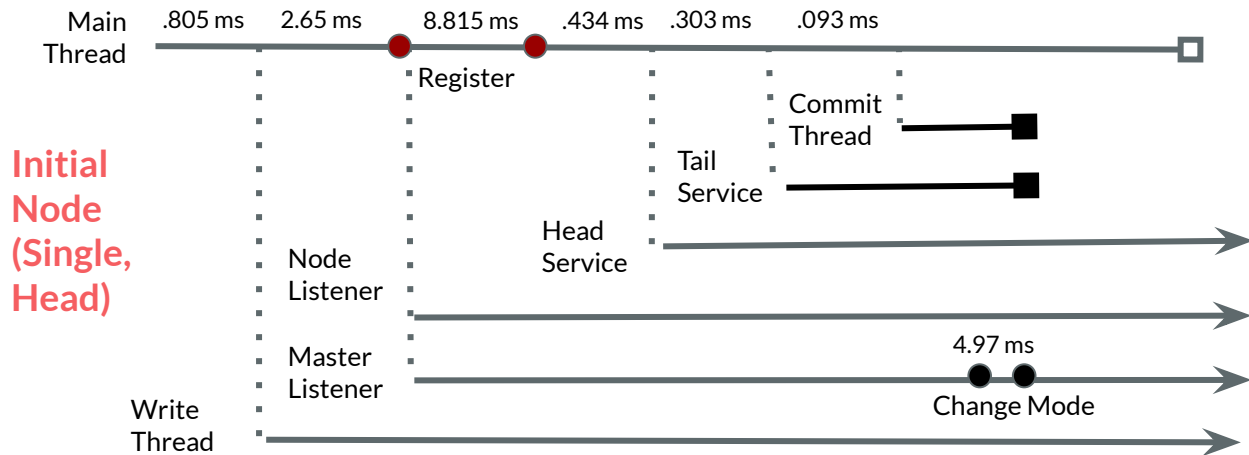
# Registration & Restoration

- 1) At start of restoration, tail pauses writes and write ack relays
- 2) Tail uses last seq# from new node to get all writes past this point from the volume, along with a copy of the replay log and sent list
- 3) Tail sends data from each of these to the new node bringing its state in line with the chain, and then resumes normal operations


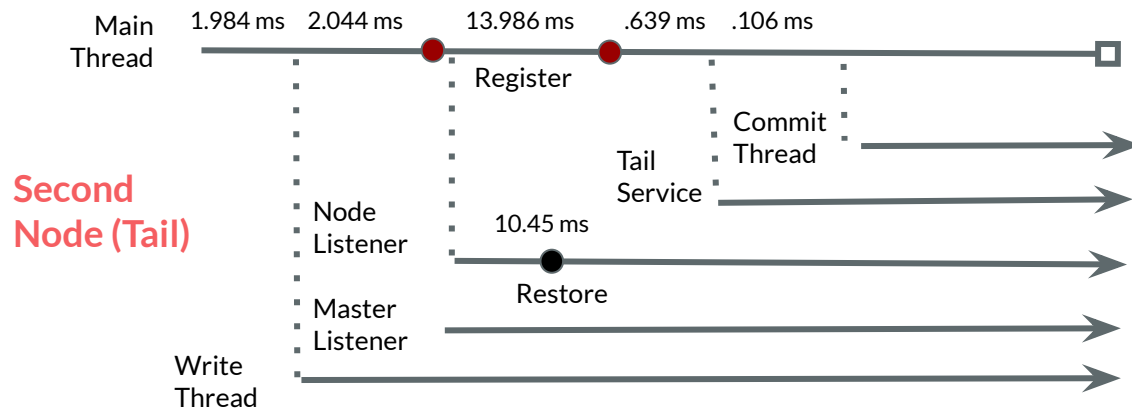
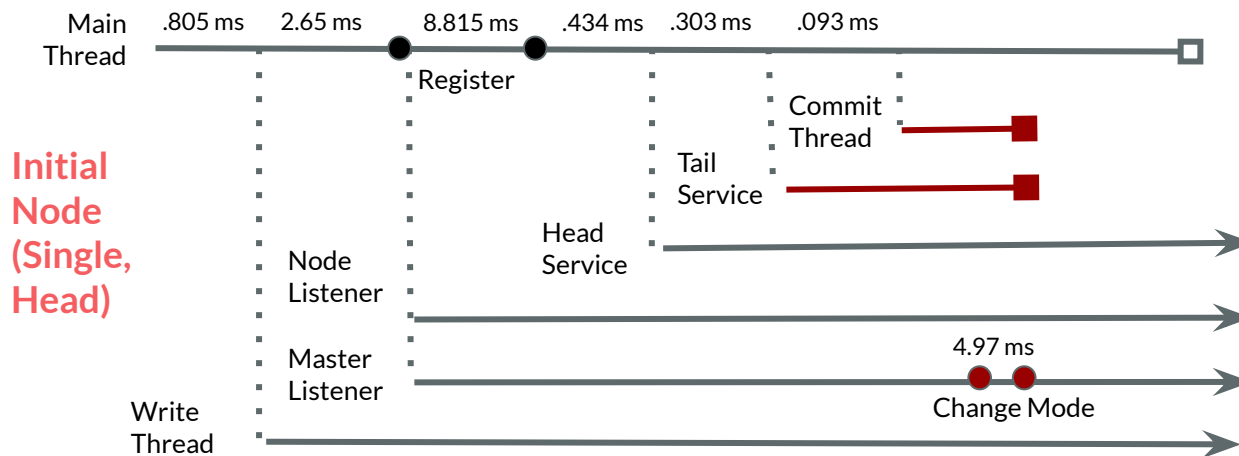





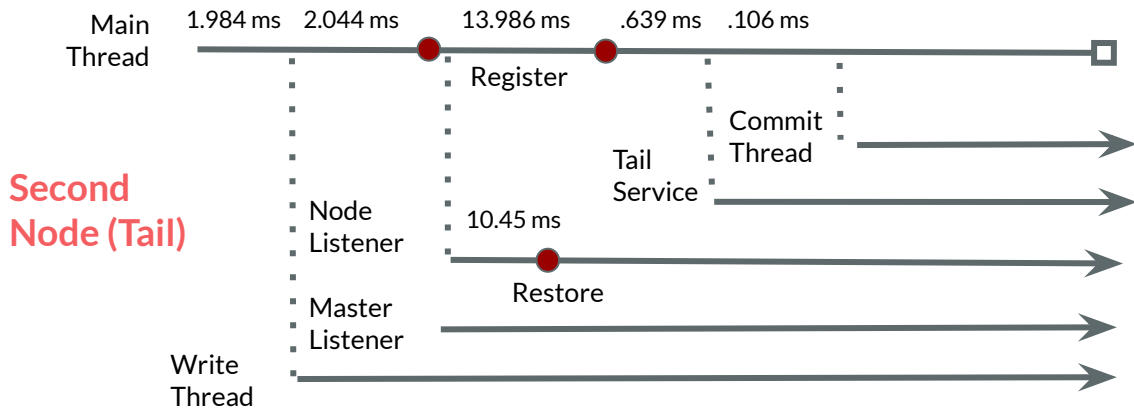
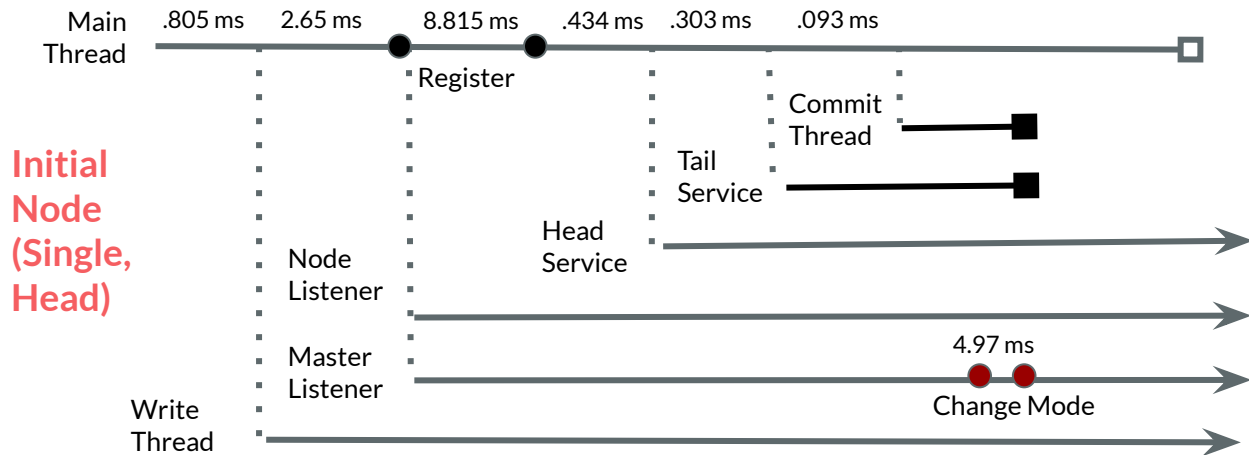
# Integration Timeline



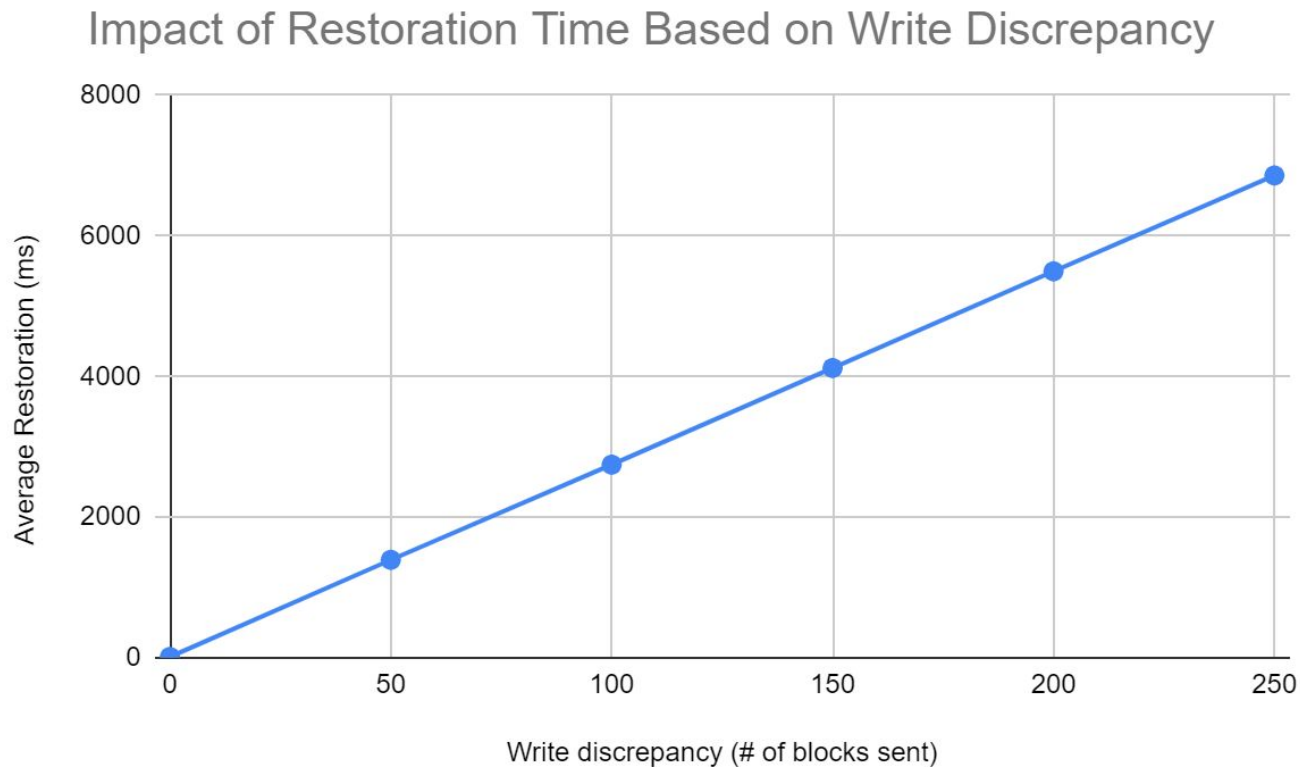
# Integration Timeline



# Integration Timeline




# Integration Timeline



# Heartbeat

- Thread wakes up and iterates through master node list sending heartbeats to each server
- On detection of failure master will
  - Identify replacement nodes and update virtual node states
  - Removes failed nodes from node list
  - Informs impacted nodes of new state and changes in up\downstream ip's
  - Update head\tail pointers as needed
- Impacted nodes make necessary modifications to recover including updating local state
- Thread goes to sleep for 5 seconds

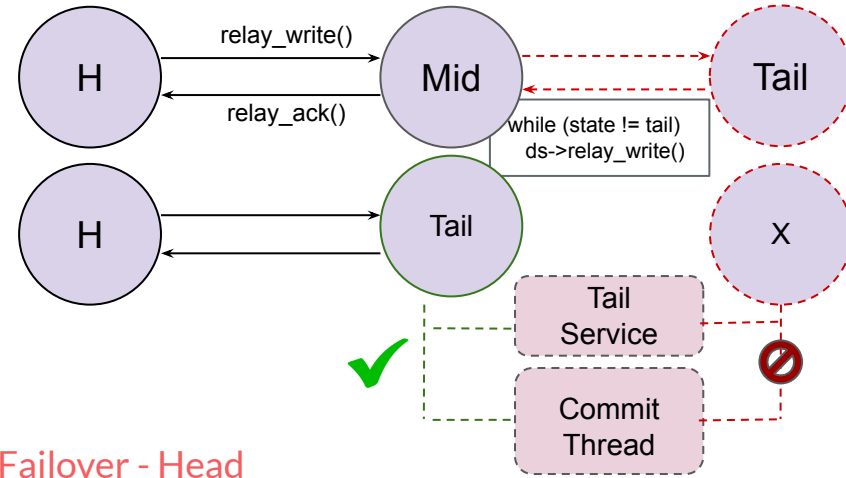
Can deal with multiple failures at once and complete chain failures



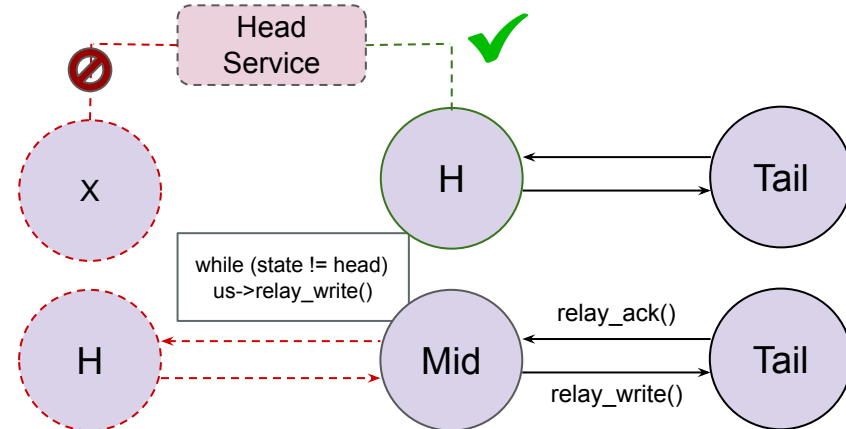
# Failover - Head\Tail

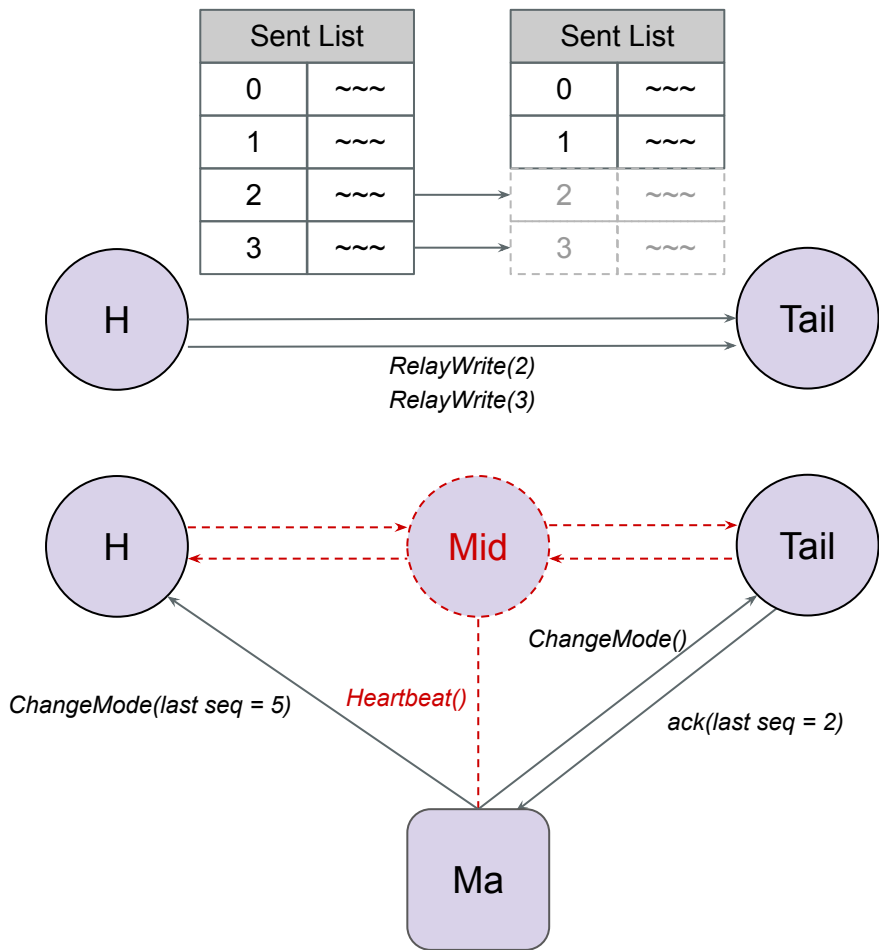
- Simple case - Head and tail failover follow similar mechanisms and have minimal impact.
- Both write and commit threads attempt to send messages while wrapped in state checks. If adjacent node fails will continue to try and send until state is updated
- New endpoint will launch relative endpoint services
- For tail failure, new tail launches relay\_write\_ack thread which starts committing sent

## Failover - Tail



## Failover - Head





## Failover - Mid


- Mid failure involves internode communication to resolve
- M+1 is notified of change by master, and responds with highest sequence number in sent list
- Master notifies M-1 of change, and send along highest sequence seen by M+1
- M-1 resends writes to M+1 that were lost during the failure

# Consistency


- Crash on command in code
  - 'Crash offset' code consisting of target IP and crash point
  - Server code decodes offset, crashes if IP/crash point match, passes on as offset 0 if not
  - Trigger different crashes at specific intermediate servers and arbitrary points by modifying client request, not server code
- Checksums: verify volume checksums along the entire chain
- Demos
  - Head failure
  - Mid failure
  - Tail failure
  - Re-integrating failed head as new tail




# Head Failure




# Mid Failure



# Tail Failure



# Re-Integration



FIN